

This repository | Search

Pull requests Issues Gist



yoonkim / CNN_sentence

Watch 54

★ Unstar 770

🔗 Fork 428

<> Code

🔔 Issues 26

🔗 Pull requests 4

📁 Projects 0

📖 Wiki

🔊 Pulse

📊 Graphs

Branch: master

CNN_sentence / conv_net_classes.py

Find file

Copy path

yoonkim init push

4abc8df on Dec 4, 2014

1 contributor

419 lines (349 sloc) 17.1 KB

Raw

Blame

History



```

1  """
2  Sample code for
3  Convolutional Neural Networks for Sentence Classification
4  http://arxiv.org/pdf/1408.5882v2.pdf
5
6  Much of the code is modified from
7  - deeplearning.net (for ConvNet classes)
8  - https://github.com/mdenil/dropout (for dropout)
9  - https://groups.google.com/forum/#!topic/pylearn-dev/3QbKtCumAW4 (for Adadelta)
10 """
11
12 import numpy
13 import theano.tensor.shared_randomstreams
14 import theano
15 import theano.tensor as T
16 from theano.tensor.signal import downsample
17 from theano.tensor.nnet import conv
18
19 def ReLU(x):
20     y = T.maximum(0.0, x)
21     return(y)
22 def Sigmoid(x):
23     y = T.nnet.sigmoid(x)
24     return(y)
25 def Tanh(x):
26     y = T.tanh(x)
27     return(y)
28 def Iden(x):
29     y = x
30     return(y)
31
32 class HiddenLayer(object):
33     """
34     Class for HiddenLayer
35     """
36     def __init__(self, rng, input, n_in, n_out, activation, W=None, b=None,
37                  use_bias=False):
38
39         self.input = input
40         self.activation = activation
41
42         if W is None:
43             if activation.func_name == "ReLU":
44                 W_values = numpy.asarray(0.01 * rng.standard_normal(size=(n_in, n_out)), dtype=theano.config.floatX)
45             else:
46                 W_values = numpy.asarray(rng.uniform(low=-numpy.sqrt(6. / (n_in + n_out)), high=numpy.sqrt(6. / (n_in + n_out)),
47                                                    size=(n_in, n_out)), dtype=theano.config.floatX)
48             W = theano.shared(value=W_values, name='W')
49         if b is None:

```

```

50         b_values = numpy.zeros((n_out,), dtype=theano.config.floatX)
51         b = theano.shared(value=b_values, name='b')
52
53         self.W = W
54         self.b = b
55
56         if use_bias:
57             lin_output = T.dot(input, self.W) + self.b
58         else:
59             lin_output = T.dot(input, self.W)
60
61         self.output = (lin_output if activation is None else activation(lin_output))
62
63         # parameters of the model
64         if use_bias:
65             self.params = [self.W, self.b]
66         else:
67             self.params = [self.W]
68
69     def _dropout_from_layer(rng, layer, p):
70         """p is the probability of dropping a unit
71         """
72         srng = theano.tensor.shared_randomstreams.RandomStreams(rng.randint(999999))
73         # p=1-p because 1's indicate keep and p is prob of dropping
74         mask = srng.binomial(n=1, p=1-p, size=layer.shape)
75         # The cast is important because
76         # int * float32 = float64 which pulls things off the gpu
77         output = layer * T.cast(mask, theano.config.floatX)
78         return output
79
80     class DropoutHiddenLayer(HiddenLayer):
81         def __init__(self, rng, input, n_in, n_out,
82                     activation, dropout_rate, use_bias, W=None, b=None):
83             super(DropoutHiddenLayer, self).__init__(
84                 rng=rng, input=input, n_in=n_in, n_out=n_out, W=W, b=b,
85                 activation=activation, use_bias=use_bias)
86
87             self.output = _dropout_from_layer(rng, self.output, p=dropout_rate)
88
89     class MLPDropout(object):
90         """A multilayer perceptron with dropout"""
91         def __init__(self, rng, input, layer_sizes, dropout_rates, activations, use_bias=True):
92
93             #rectified_linear_activation = lambda x: T.maximum(0.0, x)
94
95             # Set up all the hidden layers
96             self.weight_matrix_sizes = zip(layer_sizes, layer_sizes[1:])
97             self.layers = []
98             self.dropout_layers = []
99             self.activations = activations
100             next_layer_input = input
101             #first_layer = True
102             # dropout the input
103             next_dropout_layer_input = _dropout_from_layer(rng, input, p=dropout_rates[0])
104             layer_counter = 0
105             for n_in, n_out in self.weight_matrix_sizes[:-1]:
106                 next_dropout_layer = DropoutHiddenLayer(rng=rng,
107                                                         input=next_dropout_layer_input,
108                                                         activation=activations[layer_counter],
109                                                         n_in=n_in, n_out=n_out, use_bias=use_bias,
110                                                         dropout_rate=dropout_rates[layer_counter])
111                 self.dropout_layers.append(next_dropout_layer)
112                 next_dropout_layer_input = next_dropout_layer.output
113
114             # Reuse the parameters from the dropout layer here, in a different
115             # path through the graph.
116             next_layer = HiddenLayer(rng=rng,

```

```

117         input=next_layer_input,
118         activation=activations[layer_counter],
119         # scale the weight matrix W with (1-p)
120         W=next_dropout_layer.W * (1 - dropout_rates[layer_counter]),
121         b=next_dropout_layer.b,
122         n_in=n_in, n_out=n_out,
123         use_bias=use_bias)
124     self.layers.append(next_layer)
125     next_layer_input = next_layer.output
126     #first_layer = False
127     layer_counter += 1
128
129     # Set up the output layer
130     n_in, n_out = self.weight_matrix_sizes[-1]
131     dropout_output_layer = LogisticRegression(
132         input=next_dropout_layer_input,
133         n_in=n_in, n_out=n_out)
134     self.dropout_layers.append(dropout_output_layer)
135
136     # Again, reuse paramters in the dropout output.
137     output_layer = LogisticRegression(
138         input=next_layer_input,
139         # scale the weight matrix W with (1-p)
140         W=dropout_output_layer.W * (1 - dropout_rates[-1]),
141         b=dropout_output_layer.b,
142         n_in=n_in, n_out=n_out)
143     self.layers.append(output_layer)
144
145     # Use the negative log likelihood of the logistic regression layer as
146     # the objective.
147     self.dropout_negative_log_likelihood = self.dropout_layers[-1].negative_log_likelihood
148     self.dropout_errors = self.dropout_layers[-1].errors
149
150     self.negative_log_likelihood = self.layers[-1].negative_log_likelihood
151     self.errors = self.layers[-1].errors
152
153     # Grab all the parameters together.
154     self.params = [ param for layer in self.dropout_layers for param in layer.params ]
155
156     def predict(self, new_data):
157         next_layer_input = new_data
158         for i,layer in enumerate(self.layers):
159             if i<len(self.layers)-1:
160                 next_layer_input = self.activations[i](T.dot(next_layer_input,layer.W) + layer.b)
161             else:
162                 p_y_given_x = T.nnet.softmax(T.dot(next_layer_input, layer.W) + layer.b)
163                 y_pred = T.argmax(p_y_given_x, axis=1)
164                 return y_pred
165
166     def predict_p(self, new_data):
167         next_layer_input = new_data
168         for i,layer in enumerate(self.layers):
169             if i<len(self.layers)-1:
170                 next_layer_input = self.activations[i](T.dot(next_layer_input,layer.W) + layer.b)
171             else:
172                 p_y_given_x = T.nnet.softmax(T.dot(next_layer_input, layer.W) + layer.b)
173                 return p_y_given_x
174
175     class MLP(object):
176         """Multi-Layer Perceptron Class
177
178         A multilayer perceptron is a feedforward artificial neural network model
179         that has one layer or more of hidden units and nonlinear activations.
180         Intermediate layers usually have as activation function tanh or the
181         sigmoid function (defined here by a ``HiddenLayer`` class) while the
182         top layer is a softmax layer (defined here by a ``LogisticRegression``

```

```

183     class).
184     """
185
186     def __init__(self, rng, input, n_in, n_hidden, n_out):
187         """Initialize the parameters for the multilayer perceptron
188
189         :type rng: numpy.random.RandomState
190         :param rng: a random number generator used to initialize weights
191
192         :type input: theano.tensor.TensorType
193         :param input: symbolic variable that describes the input of the
194             architecture (one minibatch)
195
196         :type n_in: int
197         :param n_in: number of input units, the dimension of the space in
198             which the datapoints lie
199
200         :type n_hidden: int
201         :param n_hidden: number of hidden units
202
203         :type n_out: int
204         :param n_out: number of output units, the dimension of the space in
205             which the labels lie
206
207         """
208
209         # Since we are dealing with a one hidden layer MLP, this will translate
210         # into a HiddenLayer with a tanh activation function connected to the
211         # LogisticRegression layer; the activation function can be replaced by
212         # sigmoid or any other nonlinear function
213         self.hiddenLayer = HiddenLayer(rng=rng, input=input,
214                                         n_in=n_in, n_out=n_hidden,
215                                         activation=T.tanh)
216
217         # The logistic regression layer gets as input the hidden units
218         # of the hidden layer
219         self.logRegressionLayer = LogisticRegression(
220             input=self.hiddenLayer.output,
221             n_in=n_hidden,
222             n_out=n_out)
223
224         # L1 norm ; one regularization option is to enforce L1 norm to
225         # be small
226
227         # negative log likelihood of the MLP is given by the negative
228         # log likelihood of the output of the model, computed in the
229         # logistic regression layer
230         self.negative_log_likelihood = self.logRegressionLayer.negative_log_likelihood
231         # same holds for the function computing the number of errors
232         self.errors = self.logRegressionLayer.errors
233
234         # the parameters of the model are the parameters of the two layer it is
235         # made out of
236         self.params = self.hiddenLayer.params + self.logRegressionLayer.params
237
238     class LogisticRegression(object):
239         """Multi-class Logistic Regression Class
240
241         The logistic regression is fully described by a weight matrix :math:`W`
242         and bias vector :math:`b`. Classification is done by projecting data
243         points onto a set of hyperplanes, the distance to which is used to
244         determine a class membership probability.
245         """
246
247         def __init__(self, input, n_in, n_out, W=None, b=None):
248             """ Initialize the parameters of the logistic regression
249

```

```

250 :type input: theano.tensor.TensorType
251 :param input: symbolic variable that describes the input of the
252 architecture (one minibatch)
253
254 :type n_in: int
255 :param n_in: number of input units, the dimension of the space in
256 which the datapoints lie
257
258 :type n_out: int
259 :param n_out: number of output units, the dimension of the space in
260 which the labels lie
261
262 """
263
264 # initialize with 0 the weights W as a matrix of shape (n_in, n_out)
265 if W is None:
266     self.W = theano.shared(
267         value=numpy.zeros((n_in, n_out), dtype=theano.config.floatX),
268         name='W')
269 else:
270     self.W = W
271
272 # initialize the biases b as a vector of n_out 0s
273 if b is None:
274     self.b = theano.shared(
275         value=numpy.zeros((n_out,), dtype=theano.config.floatX),
276         name='b')
277 else:
278     self.b = b
279
280 # compute vector of class-membership probabilities in symbolic form
281 self.p_y_given_x = T.nnet.softmax(T.dot(input, self.W) + self.b)
282
283 # compute prediction as class whose probability is maximal in
284 # symbolic form
285 self.y_pred = T.argmax(self.p_y_given_x, axis=1)
286
287 # parameters of the model
288 self.params = [self.W, self.b]
289
290 def negative_log_likelihood(self, y):
291     """Return the mean of the negative log-likelihood of the prediction
292     of this model under a given target distribution.
293
294     .. math::
295
296     \frac{1}{|\mathcal{D}|} \mathcal{L}(\theta=\{W,b\}, \mathcal{D}) =
297     \frac{1}{|\mathcal{D}|} \sum_{i=0}^{|\mathcal{D}|} \log(P(Y=y^{(i)}|x^{(i)}, W,b)) \ll
298     \ell(\theta=\{W,b\}, \mathcal{D})
299
300 :type y: theano.tensor.TensorType
301 :param y: corresponds to a vector that gives for each example the
302 correct label
303
304 Note: we use the mean instead of the sum so that
305 the learning rate is less dependent on the batch size
306 """
307 # y.shape[0] is (symbolically) the number of rows in y, i.e.,
308 # number of examples (call it n) in the minibatch
309 # T.arange(y.shape[0]) is a symbolic vector which will contain
310 # [0,1,2,... n-1] T.log(self.p_y_given_x) is a matrix of
311 # Log-Probabilities (call it LP) with one row per example and
312 # one column per class LP[T.arange(y.shape[0]),y] is a vector
313
314 # v containing [LP[0,y[0]], LP[1,y[1]], LP[2,y[2]], ...,
315 # LP[n-1,y[n-1]]] and T.mean(LP[T.arange(y.shape[0]),y]) is
316 # the mean (across minibatch examples) of the elements in v,

```

```

316         # i.e., the mean log-likelihood across the minibatch.
317         return -T.mean(T.log(self.p_y_given_x)[T.arange(y.shape[0]), y])
318
319     def errors(self, y):
320         """Return a float representing the number of errors in the minibatch ;
321         zero one loss over the size of the minibatch
322
323         :type y: theano.tensor.TensorType
324         :param y: corresponds to a vector that gives for each example the
325         correct label
326         """
327
328         # check if y has same dimension of y_pred
329         if y.ndim != self.y_pred.ndim:
330             raise TypeError('y should have the same shape as self.y_pred',
331                             ('y', target.type, 'y_pred', self.y_pred.type))
332         # check if y is of the correct datatype
333         if y.dtype.startswith('int'):
334             # the T.neq operator returns a vector of 0s and 1s, where 1
335             # represents a mistake in prediction
336             return T.mean(T.neq(self.y_pred, y))
337         else:
338             raise NotImplementedError()
339
340     class LeNetConvPoolLayer(object):
341         """Pool Layer of a convolutional network """
342
343     def __init__(self, rng, input, filter_shape, image_shape, poolsize=(2, 2), non_linear="tanh"):
344         """
345         Allocate a LeNetConvPoolLayer with shared variable internal parameters.
346
347         :type rng: numpy.random.RandomState
348         :param rng: a random number generator used to initialize weights
349
350         :type input: theano.tensor.dtensor4
351         :param input: symbolic image tensor, of shape image_shape
352
353         :type filter_shape: tuple or list of length 4
354         :param filter_shape: (number of filters, num input feature maps,
355                             filter height,filter width)
356
357         :type image_shape: tuple or list of length 4
358         :param image_shape: (batch size, num input feature maps,
359                             image height, image width)
360
361         :type poolsize: tuple or list of length 2
362         :param poolsize: the downsampling (pooling) factor (#rows,#cols)
363         """
364
365         assert image_shape[1] == filter_shape[1]
366         self.input = input
367         self.filter_shape = filter_shape
368         self.image_shape = image_shape
369         self.poolsize = poolsize
370         self.non_linear = non_linear
371
372         # there are "num input feature maps * filter height * filter width"
373         # inputs to each hidden unit
374         fan_in = numpy.prod(filter_shape[1:])
375         # each unit in the lower layer receives a gradient from:
376         # "num output feature maps * filter height * filter width" /
377         # pooling size
378         fan_out = (filter_shape[0] * numpy.prod(filter_shape[2:]) / numpy.prod(poolsize))
379         # initialize weights with random weights
380
381         if self.non_linear=="none" or self.non_linear=="relu":
382             self.W = theano.shared(numpy.asarray(rng.uniform(low=-0.01,high=0.01,size=filter_shape),
383                                                 dtype=theano.config.floatX),borrow=True,name="W_conv")

```

```

382         else:
383             W_bound = numpy.sqrt(6. / (fan_in + fan_out))
384             self.W = theano.shared(numpy.asarray(rng.uniform(low=-W_bound, high=W_bound, size=filter_shape),
385                 dtype=theano.config.floatX), borrow=True, name="W_conv")
386             b_values = numpy.zeros((filter_shape[0],), dtype=theano.config.floatX)
387             self.b = theano.shared(value=b_values, borrow=True, name="b_conv")
388
389             # convolve input feature maps with filters
390             conv_out = conv.conv2d(input=input, filters=self.W, filter_shape=self.filter_shape, image_shape=self.image_shape)
391             if self.non_linear=="tanh":
392                 conv_out_tanh = T.tanh(conv_out + self.b.dimshuffle('x', 0, 'x', 'x'))
393                 self.output = downsample.max_pool_2d(input=conv_out_tanh, ds=self.poolsize, ignore_border=True)
394             elif self.non_linear=="relu":
395                 conv_out_tanh = ReLU(conv_out + self.b.dimshuffle('x', 0, 'x', 'x'))
396                 self.output = downsample.max_pool_2d(input=conv_out_tanh, ds=self.poolsize, ignore_border=True)
397             else:
398                 pooled_out = downsample.max_pool_2d(input=conv_out, ds=self.poolsize, ignore_border=True)
399                 self.output = pooled_out + self.b.dimshuffle('x', 0, 'x', 'x')
400             self.params = [self.W, self.b]
401
402     def predict(self, new_data, batch_size):
403         """
404         predict for new data
405         """
406         img_shape = (batch_size, 1, self.image_shape[2], self.image_shape[3])
407         conv_out = conv.conv2d(input=new_data, filters=self.W, filter_shape=self.filter_shape, image_shape=img_shape)
408         if self.non_linear=="tanh":
409             conv_out_tanh = T.tanh(conv_out + self.b.dimshuffle('x', 0, 'x', 'x'))
410             output = downsample.max_pool_2d(input=conv_out_tanh, ds=self.poolsize, ignore_border=True)
411         if self.non_linear=="relu":
412             conv_out_tanh = ReLU(conv_out + self.b.dimshuffle('x', 0, 'x', 'x'))
413             output = downsample.max_pool_2d(input=conv_out_tanh, ds=self.poolsize, ignore_border=True)
414         else:
415             pooled_out = downsample.max_pool_2d(input=conv_out, ds=self.poolsize, ignore_border=True)
416             output = pooled_out + self.b.dimshuffle('x', 0, 'x', 'x')
417         return output
418

```

